The scaling factor in self-attention

William Chuang

September 26, 2023

1 Introduction

In the context of deep learning neural networks, attention algorithms are designed for determining connections by weights between two elements in a sequential inputs. Self-attention was proposed by Vaswani et. al.[2] in 2017. Suppose a sequential input $X \in \mathbb{R}^N \times \mathbb{R}^{d_x}$ is given. Let $Q := XW_q \in \mathbb{R}^N \times \mathbb{R}^{d_q}, K := XW_k \in \mathbb{R}^N \times \mathbb{R}^{d_k}$, and $V := XW_v \in \mathbb{R}^N \times \mathbb{R}^{d_v}$ where W_k, W_v, and W_q are initiated such that (i) $W_q \in \mathbb{R}^{d_x} \times \mathbb{R}^{d_q}, W_k \in \mathbb{R}^{d_x} \times \mathbb{R}^{d_k}$, and $W_v \in \mathbb{R}^{d_x} \times \mathbb{R}^{d_v}$, (ii) the variance of the inner product of each row vector in X with each column vector in either W_q, W_k, or W_v to be unit, (iii) each row vector in Q denoted by q_i, and each row vector in K denoted by k_i to be random variables with zero means, i.e. $\mathbb{E}q_i = 0 = \mathbb{E}k_i, \forall i$, and unit variances, i.e $\mathrm{Var}q_i = 1 = \mathrm{Var}k_i$, where $i \in \{1, ..., d_k\}$, and (iv) q_i and k_j are all independent in the following ways: $\mathrm{Cov}\,(q_i, k_i) = 0, \forall i$, $\mathrm{Cov}\,(q_i, q_j) = 0, \forall i \neq j$, and $\mathrm{Cov}\,(k_i, k_j) = 0, \forall i \neq j$. By convention, let $d_k = d_q$. Then the each scalar product of column vectors in Q and K has variance d_k (see Appendix).

Definition 1 (Self-attention). *The self-attention is the following function*

$$Attention\,(Q, K, V) := \left(softmax \left(\frac{(QK^T)_i}{\sqrt{d_k}} \right)_{i=1}^N \right)_{N \times N} \circ V_{N \times d_v}.$$

Remark: The Softmax function is defined using the Boltzmann distribution function: Given a

finite sequence $\{z_i\}$, $j \in \{1, ..., n\}$, then

$$\text{Softmax}(z_i) := \frac{\exp(z_i)}{\sum\limits_j \exp(z_j)}.$$

Since the sum of the above function adds up to one, i.e. $\sum\limits_{i=1}^{n} \text{Softmax}(z_i) = 1$, it is a probability distribution. In physics, this distribution is called Boltzmann distribution. Furthermore, this distribution can be scaled by using a parameter β, and β is the inverse of the temperature of the physical system described by the distribution:

$$\text{Softmax}(\beta z_i) := \frac{\exp(\beta z_i)}{\sum\limits_j \exp(\beta z_j)}.$$

In self-attention, this β is $\frac{1}{\sqrt{d_k}}$, and T is the standard deviation of the scalar product $\sqrt{d_k}$.

Apply Softmax to column vectors in the matrix $\frac{QK^T}{\sqrt{d_k}}$, meaning each entry of the matrix $\text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)_{N \times N}$ is obtained by taking the scalar product on the i-th row of Q to multiply the $j-th$ column of K^T, then divide the scalar product by $\sqrt{d_k}$, and then apply Softmax function on each column of the matrix $\left(\frac{QK^T}{\sqrt{d_k}}\right)_{N \times N}$. Finally, take matrix multiplication between $\text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)_{N \times N}$ and $V_{N \times d_v}$.

Hence, in self-attention, each attention, i.e. weight, between each element in the sequential input to the other element in the same sequential input was weighted by the Boltzmann distribution $\text{softmax}\left(\frac{(QK^T)_i}{\sqrt{d_k}}\right)_{i=1}^{N}$.

Furthermore, because of the random initiation of each weight matrices, if this self-attention, i.e. weights assigning process were run in parallel, then the model can learn from more than one different perspectives. In [2], this notion is defined as multi-head attention. Each self-attention function $\text{Attention}(Q, K, V)_i$ is called head_i, where i is from 1 to H. The output of all H heads are concatenated to a tensor denoted by $\text{Concat}(\text{head}_1, ..., \text{head}_H)$, then defined $\text{Multihead(Q,K,V)} := \text{Concat}(\text{head}_1, ..., \text{head}_H) \circ W^O$, where $W^O \in \mathbb{R}^{H \cdot d_v} \times \mathbb{R}^{d_{\text{out}}}$.

Let $\epsilon > 0$ be given. Because of the mathematical form of the Boltzmann distribution, if the input is $z_i + \epsilon$, then $\text{Softmax}(\beta z_i + \epsilon) = \text{Softmax}(\beta z_i)$. However, a naive but natural question to ask is what if the scaling factor β in Boltzmann distribution, or the self-attention function, is modified

into the form β^ϵ or $\beta \cdot \epsilon$? It is necessary to ask this question, since we want to be able to justify why and when it makes sense to use the scaling factor $\beta = \frac{1}{\sqrt{d_k}}$ in the self-attention function. Is it sufficient to ask this question? That is the question this paper is going to answer. The next section provides a solution to this question.

2   Algorithm

To be able to make sense to use the scaling factor $\beta = \frac{1}{\sqrt{d_k}}$ in self-attention function, the four assumptions are necessary.

Within the four assumptions, the first assumption is usually fixed when self-attention is considered to be applied.

Since the way query and key matrices are defined, i.e. $Q = XW_q$ and $K = XW_k$, assumption (iii) and (iv) are based on assumption (ii).

Suppose the assumptions (ii) is relaxed, the variance of the scalar product $q_i k_i$ is approximated to $\left(d_x \sigma_X^2 \sigma_{W_q}^2 \right) \cdot \left(d_x \sigma_X^2 \sigma_{W_k}^2 \right) \cdot d_k$. Since the training set is given, i.e. it is fixed, and so are d_x and d_k, σ_{W_q} and σ_{W_k} are the actual moving factors.

What can cause σ_{W_q} and σ_{W_k} to be changed? In the training phase, the gradient descent algorithm or the backpropagation is used and the algorithm is going to shift the initial probability distribution of each entry in the random matrices W_q and W_k to the final probability distribution at a local minimum in the moduli space (or on the loss surface).

Hence, the scaling factor $\beta = \frac{1}{\sqrt{d_k}}$ is only valid at the initial point given the four assumptions at the initial point in the moduli space, and it can hold true if the thermodynamic system described by the Boltzmann distribution is in statics, i.e. its temperature is fixed. However, since the initial point is random and unlikely to be the final point in the moduli space throughout the training mode, the assumption (ii) is not valid right after the training of the model that used self-attention function is initiated. Without assumption (ii), assumption (iii) and (iv) are invalid throughout the training and predicting modes.

To solve this problem, the solution could be written as an algorithm that it measures σ_{W_q} and σ_{W_k} and updates the scaling factor β in after each iteration or batch during the training and the

best scaling factor β that is going to be used in the predicting mode could be determined by the best model chosen from the ensemble of all trained models.

This also answers the question asked in the previous section that it shows the sufficiency to ask the question.

In practice, 709 is the upper bound of the exponent of the exponential function based e in lots of machines before the machines start to overflow. Likewise, on each machine, there exists a non-negative integer N such that the machine start to overflow the result. Hence, an alternative algorithm is: firstly try $\beta \in \left\{ 10^N, \frac{1}{10^N} \right\}$ for $N = 0, 1, 2, ...$, then run a binary search to narrow down the range of the best scaling factor β during the training after either each iteration or batch.

A further meaning of this algorithm is showed in the next section by using an empirical example.

3   Empirical example

The task of the empirical example is to flip a given string of non-negative integers. In other words, the goal is to train a multihead in a transformer to learn a function that map each index i of an integer in a string to the index $(n - i) + 1$ in the new string that will be return by the model. Since the model needs to learn to swap the first and last elements in the string, it could be a challenging task for recurrent neural nets if the length of the input string is large.

In this empirical example, when the string length is 20, the range of the non-negative integers is between 0 and 100, and the number of samples used in training is 15000, 1000 samples for validation, and 100000 samples for testing.

Then, without running the above alternative algorithm, i.e. when the scaling factor β is $\frac{1}{\sqrt{d_k}}$ the validation accuracy is 1.51%, and the test accuracy is 1.50%. To have a stable test accuracy above 95%, the training sample size needs to be increased from 15000 to 25000.

However, by applying the above alternative algorithm, the scaling factor is in $O(1)$, a constant ~ 2.2, with training sample size at 15000, the validation accuracy and test accuracy can already reach above 95% stably. Thus, the above algorithms can not only compress the training set when the training samples are limited or rare, but also reduce the training time to reach the best accuracy.

Furthermore, the above the model that achieved the above accuracy did not achieve it by

memorizing the training examples. Since a 100% accuracy is also achievable within the ensemble, and once achieve that 100% accuracy the model can flip any input string with the same length and the range of the values of its each word. However, if it was by memorizing, then the model had to memorize $100^2 0$ strings with length 20 which is largely exceed the number of parameters of the model (less than 8000).

4 Discussion

5 Conclusion

6 Appendix: Deriving the scaling factor $\frac{1}{\sqrt{d_k}}$ used in [2]

Let X and Y be random variables.

Firstly, recall the definition of variance[1]:

Definition 2.
$$\sigma^2(X) := \inf_{a \in \mathbb{R}} \mathbb{E}\left[(X - a)^2\right].$$

Hence if $X \notin L^2$, then $\mathrm{Var}[X] = \infty$. If $X \in L^2$, then $\mathbb{E}\left[(X - a)^2\right] = \mathbb{E}(X^2) - 2a\mathbb{E}(X) + a^2$ has a minimum at $a = \mathbb{E}(X)$, and the expression of the variance of X can be rewritten using the minimum:
$$\mathrm{Var}[X] = \mathbb{E}\left[(X - a)^2\right] = \mathbb{E}\left[(X - \mathbb{E}[X])^2\right].$$

Proposition 1.

$$Var[X] = \mathbb{E}(X^2) + (\mathbb{E}[X])^2 - 2(\mathbb{E}[x])^2 = \mathbb{E}(X^2) - (\mathbb{E}X)^2.$$

Proof.
$$\mathrm{Var}[X] = \mathbb{E}\left[X^2 + (\mathbb{E}X)^2 - 2X\mathbb{E}X\right]$$

$$= \mathbb{E}(X^2) + (\mathbb{E}[X])^2 - 2(\mathbb{E}[x])^2 = \mathbb{E}(X^2) - (\mathbb{E}X)^2.$$

$\square$

Secondly,

Definition 3. *the covariance of X and Y is:*

$$Cov[X, Y] := \mathbb{E}\left[(X - \mathbb{E}X)(Y - \mathbb{E}Y)\right].$$

Proposition 2.

$$Cov[X, Y] = \mathbb{E}XY - \mathbb{E}X\mathbb{E}Y.$$

Proof.

$$\text{Cov}[X, Y] = \mathbb{E}\left[XY - X\mathbb{E}Y - Y\mathbb{E}X + \mathbb{E}X\mathbb{E}Y\right]$$

$$= \mathbb{E}XY - \mathbb{E}X\mathbb{E}Y - \mathbb{E}X\mathbb{E}Y + \mathbb{E}X\mathbb{E}Y$$

$$= \mathbb{E}XY - \mathbb{E}X\mathbb{E}Y.$$

$\square$

Then, the expectation value of X^2Y^2 can be written in the following expression:

$$\mathbb{E}X^2Y^2 = \text{Cov}[X^2, Y^2] + \mathbb{E}X^2\mathbb{E}Y^2$$

$$= \text{Cov}[X^2, Y^2] + (\text{Var}X + (\mathbb{E}X)^2)(\text{Var}Y + (\mathbb{E}Y)^2).$$

Proposition 3. *If X and Y are independent and both have zero means, then*

$$Var[XY] = VarX\,VarY.$$

Proof. By applying the above results, the variance of the product of X and Y can be derived:

$$\text{Var}[XY] = \mathbb{E}X^2Y^2 - (\mathbb{E}XY)^2$$

$$= \mathbb{E}X^2Y^2 - (\mathbb{E}Y\mathbb{E}X + \text{Cov}[X, Y])^2$$

$$= \operatorname{Cov}[X^2, Y^2] + (\operatorname{Var}X + (\mathbb{E}X)^2)(\operatorname{Var}Y + (\mathbb{E}Y)^2) - (\mathbb{E}Y\mathbb{E}X + \operatorname{Cov}[X, Y])^2.$$

Since X and Y are independent, then $\operatorname{Cov}[X^2, Y^2] = 0$, and

$$\operatorname{Var}[XY] = (\operatorname{Var}X + (\mathbb{E}X)^2)(\operatorname{Var}Y + (\mathbb{E}Y)^2) - (\mathbb{E}Y\mathbb{E}X + \operatorname{Cov}[X, Y])^2.$$

Since X and Y are independent and both have zero means, i.e. $\mathbb{E}X = 0 = \mathbb{E}Y$, then

$$\operatorname{Var}[XY] = \operatorname{Var}X\operatorname{Var}Y.$$

$\square$

Proposition 4. *Let $X_i, i \in \{1, ..., n\}$ be n independent variables.*

$$Var\left(\sum_i X_i\right) = \sum_i VarX_i.$$

Proof. Then

$$\operatorname{Var}\left(\sum_i X_i\right) = \mathbb{E}\left(\left(\sum_i X_i\right)^2\right) - \left(\mathbb{E}\sum_i X_i\right)^2$$

$$= \mathbb{E}\left(\sum_i \sum_j X_i X_j\right) - \left(\sum_i \mathbb{E}X_i\right)^2$$

$$= \sum_i \sum_j (\mathbb{E}X_i X_j - \mathbb{E}X_i \mathbb{E}X_j)$$

$$= \sum_i \sum_j \operatorname{Cov}(X_i, X_j).$$

Since X_i are independent, then $\operatorname{Cov}(X_i, X_j) = 0, \forall i \neq j$, and

$$\operatorname{Var}\left(\sum_i X_i\right) = \sum_i \operatorname{Cov}(X_i, X_i)$$

$$= \sum_i \operatorname{Var}X_i.$$

$\square$

Then we can derive the assumption in [2] for the variance of QK^T written in the hidden dimension d_k:

Proposition 5. *Let q_i and k_i be random variables with zero means, i.e. $\mathbb{E}q_i = 0 = \mathbb{E}k_i, \forall i$, and unit variances, i.e $\mathrm{Var}q_i = 1 = \mathrm{Var}k_i$, where $i \in \{1, ..., d_k\}$, and they are all independent in the following ways:*

$$Cov\,(q_i, k_i) = 0, \forall i,$$

$$Cov\,(q_i, q_j) = 0, \forall i \neq j, \ and$$

$$Cov\,(k_i, k_j) = 0, \forall i \neq j.$$

Then

$$Var\,\left(QK^T\right) = d_k.$$

Proof.

$$\mathrm{Var}\,\left(QK^T\right)$$

$$= \mathrm{Var}\,\left(\sum_{i=1}^{d_k} q_i k_i\right)$$

$$= \sum_i \mathrm{Var}\,(q_i k_i)$$

$$= \sum_i \mathrm{Var}\,(q_i)\,\mathrm{Var}\,(k_i)$$

$$= \sum_i 1 \cdot 1$$

$$= d_k.$$

$\square$

The first equality is because in [2], Q and K are row vectors, and QK^T is taking their scalar product. The second equality is because $\mathrm{Cov}\,(q_i, q_j) = 0, \forall i \neq j$, and $\mathrm{Cov}\,(k_i, k_j) = 0, \forall i \neq j$. Then, the variance of the sum is the sum of variances. The third equality is because of the conditions $\mathbb{E}q_i = 0 = \mathbb{E}k_i, \forall i$ and $\mathrm{Cov}\,(q_i, k_i) = 0, \forall i$, hence we have the variance of the product is the product of variances.

The authors in [2] want to scale the result of QK^T using the standard deviation of the scalar product, so QK^T is multiplied by a factor

$$\frac{1}{\sqrt{\text{Var}\left(QK^T\right)}} = \frac{1}{\sqrt{d_k}}.$$

References

[1] Gerald B Folland, *Real analysis: modern techniques and their applications*, vol. 40, John Wiley & Sons, 1999.

[2] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin, *Attention is all you need*, Advances in neural information processing systems **30** (2017).